

Daniel Suh

Software Engineer

Sacramento, CA | (916) 945-4584 | dsuh3508@gmail.com | linkedin.com/in/danielsuh8205 | github.com/dsuh02 | dsuh02.github.io

SUMMARY

I am a software engineer at Persist AI Formulations, where I own our lab information management system end to end and build the infrastructure under a multi-agent LLM product. I work primarily in PostgreSQL and FastAPI with React and TypeScript on top, and I run what I build: I operate the same databases and services locally, on on-prem Linux hosts, and in the cloud. Most of my recent work has been large-scale data migrations, retrieval over a million vectors behind measured quality gates, and production reliability.

EDUCATION & LEADERSHIP

University of Southern California

B.S. Computer Science · May 2025

CovidHacks 2020, Co-founder. Built the event site and led sponsorship outreach for an international virtual hackathon that drew 110+ participants and \$9,400 in prizes (covidhacks2020.devpost.com).

EXPERIENCE

Persist AI Formulations

Sacramento, CA

Software Engineer · Scrum Master

May 2025 - Present

- Architected the company's lab operations platform end to end: an enum-safe PostgreSQL containment hierarchy with append-only immutable history, FastAPI endpoints for movement and state replay, and a React/TypeScript timeline for inspecting live containment trees. Built the instrument layer feeding it, including OPC UA servers and clients (asyncua), serial command bridges, and protocol schedulers driving real hardware. Later took its hot path from **roughly 100 database round trips to 2** by rewriting the bulk move endpoint as one modifying CTE.
- Consolidated **21 HTTP microservices into a single FastAPI process** coordinating an orchestrator and 17 specialist LLM agents, replacing cross-service calls with in-process tool dispatch. Removed every per-request network hop and reduced the full request path to one stack trace.
- Led migration of a **71-database data plane** (40 SQLite files and 19 on-disk vector indexes) onto a single **PostgreSQL/pgvector** cluster spanning 49 schemas, 388 tables, and roughly 14.5M rows. Validated in production with the legacy disk volume fully detached: 33-second clean boot, 15 of 15 vector stores reachable, 585 of 586 upstream calls at HTTP 200, and no filesystem fallbacks.
- Rebuilt the retrieval layer end to end with a structure-aware chunker, 22 corpora re-embedded at 3,072 dimensions (adding 121K chunks), and a two-stage search path feeding approximate HNSW recall into an exact float32 re-rank. Gated the cutover on **recall@10 of at least 0.95**, which held across all 16 migrated corpora, 14 of them at 1.000.
- Hardened production reliability. Traced connection-pool exhaustion to sessions held open across SSE streams and multi-second upstream calls, audited all **26 affected handlers**, and shipped pool timeouts plus session-scoped hot paths. Separately closed six boot-validation gaps that let /health/ready return 200 over a keyless or database-dead stack. Four adversarial review passes found zero regressions.
- Designed encryption at rest for client deployments: SQLCipher for relational stores and an AES-256-GCM envelope for index, model, and document artifacts, all keyed by one master key fetched at runtime. Extended it to column-level sealing through custom SQLAlchemy TypeDecorators across **578 columns in 32 databases**, verified 578 of 578 with retrieval quality unchanged.
- Built the client factory that all upstream model traffic routes through, supporting **three independent auth modes per provider** (static key, Azure Managed Identity, and no-auth gateway passthrough) so traffic can stay inside a customer's own cloud tenant. Migrated roughly 50 call sites off raw SDK constructors to clear a 93-error production 401 cluster, and fixed a recursive-lock deadlock that froze the event loop on the first token fetch.
- On the client-facing portal, integrated a rules-based pricing engine with the correctness properties billing requires: idempotency keys, a quoted-hash guard that returns 409 rather than charging a stale price, and order creation plus ledger charge committed atomically under a row lock. Migrated a legacy inventory domain model across the React/TypeScript and FastAPI stack behind four blocking CI gates (**382 backend and 783 frontend tests**), with Playwright validation catching two live production bugs.

Software Engineer Intern

Jan 2025 - May 2025

- Cut protocol design time by roughly **60%** with automation scripts that replaced manual experiment setup.
- Ported and modernized a legacy Python web application, adding microsphere size prediction and PubChem-backed compound-similarity search to shorten R&D iteration cycles.

USC projects: Spotted (Jan - Jun 2024), geographic data visualization on Google Maps APIs over MySQL, with AJAX-driven views for reload-free interaction and Java servlets handling data and file management. **Emotify** (Aug - Dec 2023), OpenAI and Spotify API integration, cookie-based session management, servlet-rendered dynamic views, and SQL stored procedures securing the authentication flow.

SKILLS

Languages: Python, TypeScript / JavaScript, Java, C / C++, SQL, Bash

Backend & Data: FastAPI, Flask/Jinja2, SQLAlchemy (async ORM, pooling, AsyncSession), Pydantic v2, PostgreSQL, pgvector, SQLite / SQLCipher, MySQL, Alembic, schema migrations, REST API design, WebSockets, SSE, httpx, asyncio

Frontend: React, Vite, TypeScript, Zustand, Tailwind, Recharts

AI & Retrieval: multi-agent orchestration, LLM tool calling, RAG, vector search (pgvector HNSW / IVFFlat, FAISS), embedding pipelines, OpenAI / Anthropic / Google SDKs, provider routing and cost tiering, recall and eval gates

Infrastructure & Testing: Docker, Azure (Container Apps, Managed Identity, Key Vault, API Management), CI/CD (Azure DevOps), Linux, SSH, Git, Nuitka, uv, OPC UA (asyncua), pytest, Playwright, Vitest, Agile/Scrum, code review